

ASHWIN SATISH

(857)-832-0951 | ashwinsatishjapps@gmail.com | [linkedin.com/in/ashwinsatishp](https://www.linkedin.com/in/ashwinsatishp) | github.com/NamikazeAsh | [Portfolio](#)

EDUCATION

Northeastern University - Khoury College of Computer Science

Master of Science, Data Science

Machine Learning, NLP, MLOps, Algorithms, Data Mining, Database, Capstone

Sept 2024 - Dec 2026 | Boston, MA

GPA: 3.83/4.0

Christ University

Bachelor of Science, Computer Science

June 2020 - May 2023 | Bangalore, India

GPA: 3.53/4.0

EXPERIENCE

Institute for Experiential AI | Machine Learning Developer

Jan 2026 - Present | Boston, MA

- Engineered an LLM extraction pipeline in Python, mining **35M** PubMed records to identify **278K+** trial-linked publications & transforming **13K+** papers into structured data for clinical-trial knowledge graphs ([Trialsynth](#)).
- Designed & shipped a clinical-trial data model in Neo4j (**115K+** nodes, **116K+** relationships), exposing it via REST APIs & a production Flask app ([CoGEx](#)).
- Architected a custom entity-normalization pipeline mapping **45K+** clinical-trial adverse events to standard medical ontologies, raising coverage from **75.1%** to **85.4%** (+**4.6K** events) with zero regressions.
- Benchmarked LLMs on AWS Bedrock for clinical-data extraction across cost, recall, & token efficiency, cutting per-paper inference cost **~65%** while retaining **90% recall** against the baseline model.

Northeastern University | Machine Learning Researcher

Jan 2025 - Dec 2025 | Boston, MA

- Boosted true positives from **2% to 77%** (maintaining zero false positives) for Lyme cell detection via 7-model CNN ensemble.
- Built a **Python** image preprocessing pipeline for **30,000+** MATLAB cell images, saving the biology team weeks of manual work.
- Improved accuracy by **13%** via augmentation & L2/dropout regularization, achieving production-ready performance.
- Mentored & conducted office hours as a Teaching Assistant for **300+ students** in linear algebra, statistics & **NumPy** demos.

CarerAI | Data Science Intern

Jul 2023 - May 2024 | Gurgaon, India

- Prevented **2,000+ unnecessary procedures** for customers by uncovering **500+** surgery avoidance cases through KNN analysis of patient records against historical data.
- Leveraged EHR for a **risk-score**, improving classification accuracy by **26.4%**, increasing confidence in surgery avoidance cases.
- Improved KNN's score by **20%** using DBSCAN to remove record outliers, reducing false positives in surgery recommendations.

SKILLS

- Languages:** Python, SQL, R, C, Java, TypeScript, JavaScript
- Libraries & Frameworks:** Scikit-learn, TensorFlow, PyTorch, Keras, Hugging Face, OpenCV, Pandas, NumPy, PySpark
- Cloud & Development:** AWS, GCP, Azure, PostgreSQL, Snowflake, Docker, Neo4j, Bedrock, Lambda, Kubeflow
- LLM & GenAI Tooling:** Transformers, BERT, LoRA, AutoML, T5, LangChain, LangGraph, Ollama, PydanticAI, MLFlow

PROJECTS

AI-Powered Pokémon Battle Strategy Engine | Ollama, ChromaDB, Sentence-Transformers, Docker, Streamlit | [Github](#)

- Built & dockerized hybrid **Multi-Agent System (MAS)** integrating Llama 3.2 with deterministic backend for strategy analysis with a recommendation system for optimal team synergies by ingesting **1M+** aggregated usage **stats** via Python **ETL** pipeline.
- Engineered a RAG pipeline via ChromaDB & Sentence-Transformers, indexing **40,000+** strategy docs to ground LLM inference.

PantryPilot: Pantry & Recipe Assistant (2nd, MLOps Expo: Google) | LoRA, PostgreSQL, Airflow, Docker, GCP, Ollama | [Github](#)

- Contributed to the fine tuning, evaluation & automated deployment (CI/CD) on **GCP** of Llama 3.2 3B for personalized inventory recipes, resulting in **50% reduction** in dietary violations via LoRA adaptation.
- Collaborated on **DPO feedback** with Teacher model (Llama 3 70B) for preference labelling, VLM for OCR & Veo 3.1 for video.

Airlytics for Fitbit: Analysis, Forecast & Health Coach | FastAPI, SQLAlchemy, SQLite, Groq, Scikit-learn, Pandas | [Github](#)

- Built personalized readiness-forecasting models & dashboards for Fitbit Air's data, beating a naive baseline by **16%**.
- Built LLM health coach that provides health insights on retrieved data, validated by scoring tool-call accuracy & grounding rate.

arXiv Research Assistant: Agentic RAG for Q&A | OpenSearch, LangChain, LiteLLM, NetworkX, Weights & Biases | [Github](#)

- Built an agentic RAG assistant answering questions over arXiv papers via hybrid **BM25/kNN** retrieval and a **GraphRAG** graph.
- Engineered a semantic chunker & ingestion pipeline, indexing **1785** chunks from **10** arXiv papers in under **30s** (79/s).